

Feature-Reusing Local Keypoint Refinement for Monocular Pallet Pose Estimation: Frozen Estimator Applications

Internal research draft—author details pending

Abstract—Accurate image landmarks are important to monocular pallet pose measurement, but improving localization need not require changing an established detector. We study a fixed, synthetic-supervised local keypoint refiner that reuses two feature scales of a frozen RGB estimator. In the completed YOLO application, the 18,962-parameter module predicts a distribution over local corner displacements while preserving bounding boxes, confidence scores, instance selection, and the center landmark. For this application, on 319 reused development images from 13 sessions, the mean over three independently trained seeds of pooled landmark median error changes from 6.616 to 5.905 pixels. The paired-session interval for the reduction is [0.411, 1.174] pixels. Complete-denominator keypoint accuracy, error tails, geometry-reference pose errors, and desktop latency are reported separately. A matched-evidence direct regression control and a source-derived PoseFix pallet adaptation test distinct refinement alternatives. The P-prior median difference is 0.336 pixels (paired-session interval [0.063, 0.585]). The interval favors the prior, not P. This is a fixed-budget development comparison. These measurements establish the scope of the development evidence, not an unseen-session guarantee. Independent real-session confirmation is not yet available. The extension applies the same correction principle to a separately frozen DOPE estimator with new source-trained heads and explicit missing-point accounting. DOPE results remain pending in this revision. ResNet18 results remain pending in this revision. Neither a universal model-agnostic result, an original expectation operator, nor operational handling reliability is claimed.

Index Terms—Monocular vision, pallet pose, keypoint refinement, synthetic supervision, feature reuse, measurement evaluation.

I. Introduction

Pallet localization provides geometric input to inspection, tracking, and material-handling systems. A monocular RGB camera offers a compact sensing interface, but image-level detection is not the same quantity as accurate geometric measurement. The same correctly detected instance can have displaced corners; a downstream perspective-n-point (PnP) solver can amplify some image errors depending on geometry and viewing configuration. It is therefore useful to distinguish detection, conditional landmark precision, and pose measurement when evaluating a refinement method.

Synthetic training can supply image supervision without manually annotating every pallet training image. Synthetic RGB pallet pose estimation is already established,

for example through the Euro-pallet study of Knitt et al. [1]. We do not claim its invention, a measured reduction in annotation labor, or a universal synthetic-to-real solution. Instead, we ask a narrower question: can a small local module improve the residual geometry of a fixed estimator while leaving its detection outputs unchanged?

We investigate the already-trained point refiner P. It reads existing detector features and scores candidate displacements around each predicted corner. The detector is frozen, but this is not a guarantee that corrected geometry is harmless: a wrong displacement can still degrade a keypoint or a pose. Our evaluation therefore retains tails, individual sessions, and complete ground-truth denominators alongside the central error statistic. The original application depends on the YOLO feature interface. A second, separately specified DOPE application tests whether the same local readout principle remains useful with different frozen features and error patterns. It requires a new adapter and new heads, not zero-shot transfer of the YOLO weights. A finite set of fixed applications cannot establish universal model-agnostic behavior.

The third application is specified as a full-image, nine-landmark SimpleBaseline-derived ResNet18 with ImageNet initialization. Its baseline training, frozen P/D head fits and measured outcomes remain pending in this revision. Completion of DOPE alone does not satisfy the three-estimator evidence requirement.

The contributions are an explicitly specified feature-reusing refinement interface, within-estimator comparisons of the local distribution readout with a direct residual control, and an auditable protocol tying original-pixel outputs to failure-aware geometry and measured cost. Historical YOLO line-structured and prior-derived comparisons remain visible. The second-estimator claim is conditional on the separately reported DOPE results, including adverse outcomes. These are system and empirical contributions to be assessed against the evidence. We do not introduce a new expectation operator, a low-rank adaptation algorithm, or a new safety controller. Figure 1 separates the learned correction from the frozen detector and the geometry inputs.

II. Related Work and Comparison Scope

A. Pallet perception and synthetic supervision

DOPE[2] estimates semantic cuboid landmarks from RGB. Knitt et al. [1] trained a DOPE-based pallet




Fig. 1. Information paths of the completed YOLO application; the DOPE extension has its separate frozen adapter. P uses image features and predicted points/box. Camera intrinsics and registered physical dimensions enter the downstream pose solver, not P. Synthetic labels supervise P during training only. Boxes, scores, selected instance, and center are preserved; corner geometry can change.

estimator using synthetic imagery and evaluated RGB Euro-pallet pose. More recently, Miura et al. [3] studied monocular pallet pose using metric-depth predictions and ground-plane information. These works motivate the application and establish that RGB pallet estimation and synthetic data are not new in themselves. Their datasets and auxiliary information differ from ours; published numbers are not subtracted from the present development results.

B. Pose and coordinate refinement

PoseFix[4] refines an input pose using the image and a separate backbone. Its human-pose training procedure synthesizes pose errors from human-keypoint statistics. Our purpose adaptation replaces those human assumptions with frozen source-training predictions from the pallet detector. It is named PoseFix-derived pallet9, matched-exposure adaptation, not a reproduction of the original human benchmark. Integral regression[5] establishes differentiable coordinate expectation as prior art. P applies an expectation to candidate displacements; the expectation itself is not a claimed invention.

CRT-6D[6] uses object-surface keypoint features in cascaded pose refinement. It is relevant to sparse feature-based correction, but has a different object-surface and pose contract. The inspected official repository at the recorded commit contains no executable network; no substitute transformer is reported as CRT-6D. Our named prior experiment is limited to the declared PoseFix-derived family rather than expanding a benchmark after seeing results.

C. Feature reuse and small side networks

Side-tuning[7] relates to adapting a fixed representation through a separate learned path. P shares the general motivation of retaining a backbone and training a smaller addition. Nevertheless, P does not implement LoRA, and no new parameter-efficient fine-tuning benchmark is established. Both P and the direct control D have

small learned heads. Their difference therefore cannot be attributed simply to parameter efficiency. Freezing the backbone also does not make inference overhead zero.

III. Fixed Refinement Method

A. Sensor, detector, and coordinates

The historical baseline R0 is a YOLO26n pose estimator producing a box, confidence, and nine points per candidate instance. The selected instance is the highest-confidence prediction under the existing pipeline. Corner indices 0–7 retain the registered cuboid convention, and index 8 retains the established center definition. The wrapper does not use a ground-truth box or pose to select the input instance. Ground truth is used downstream for training supervision and evaluation matching only.

R0 was initialized from a pretrained YOLO26n pose checkpoint, and its stored provenance identifies COCO-pose training in that checkpoint. The pallet run also records pretrained initialization. Thus, “synthetic-supervised refinement” refers to the training of the additional module, not a claim that no real annotations ever contributed to the full estimator. Real labels are also used in development and evaluation. R0 is not retrained in the present comparison.

The historical image pipeline is preserved: rectangular detector preprocessing at the registered input size, stride-aware padding, and original-image restoration. Prepared synthetic images already include their recorded reflection padding and are not padded a second time. Let A be the diagonal image-to-network scaling matrix and $\mathbf{o}$ its offset. Original points $\mathbf{p}$ and network-input points $\tilde{\mathbf{p}}$ satisfy $\tilde{\mathbf{p}} = A\mathbf{p} + \mathbf{o}$. YOLO letterboxing gives $A = gI$; the DOPE raster can have unequal actual horizontal and vertical scales. Corrections use A^{-1} on displacements and add them to the retained original points. A zero correction therefore preserves the exact baseline coordinates.

B. Local candidate evidence

P adapts P3/P4 features with small learned channel projections. Around each valid corner, 13 directions and 17 radii define 221 displacement candidates within 0.08 predicted-box diagonal. A zero-displacement null candidate completes the 222-way distribution. Each non-null candidate reads a 4×8 stencil from both feature scales using the registered bilinear sampling geometry. A shared patch encoder and corner-role/context features feed candidate and null scorers. No explicit edge or Hough representation is required by P.

Let d_b denote the predicted box diagonal in network-input pixels, $\mathbf{v}_j$ the normalized candidate offsets, and z_{kj} the score for corner k . For the frozen temperature T and scale λ , the proposed displacement is

$$\delta_k = \lambda d_b \sum_j \frac{\exp(z_{kj}/T)}{\sum_l \exp(z_{kl}/T)} \mathbf{v}_j. \quad (1)$$




Fig. 2. The implemented candidate lattice and one actual sampling stencil in box-diagonal units. The blue stencil is shown at a fixed candidate index, not selected for performance. The probability expectation in (1) combines candidate displacements; no illustrative weights are presented as measured outputs. Geometry and source hashes are recorded in FIGURE_MANIFEST.json.

The null candidate has $\mathbf{v}_{\text{null}} = \mathbf{0}$. Given an original-image diagonal d_I and frozen cap fraction α , movement and restoration are

$$\Delta_k = A^{-1}\delta_k, \quad (2)$$

$$\bar{\Delta}_k = \Delta_k \min \left(1, \frac{\alpha d_I}{\|\Delta_k\|_2} \right), \quad (3)$$

$$\mathbf{p}'_k = \mathbf{p}_k + \bar{\Delta}_k, \quad (4)$$

with unit scaling for a zero movement and no clipping when the selected cap is absent. Historical YOLO P retains its selected $\lambda = 1$, $\alpha = 0.01$, and per-seed temperatures. DOPE uses the same source-only selection grid; its selected values are supplied with its own result records. Unsupported points are unchanged; point 8 is copied exactly. Detection boxes, confidences, candidate order, and nonselected points are preserved. These contracts do not imply unchanged or always improved 6D geometry.

C. Training and internal controls

P learns a soft distribution over candidate displacements from synthetic corner residuals. The finite target support and per-frame reduction are retained from the fixed implementation. The historical YOLO P, D, and L fits each use their recorded three seeds and fixed exposure budget; those models are not retrained. YOLO P has 18,962 learned parameters. D reuses the same sampling and shared feature descriptors, but regresses normalized corner residuals with an L1 objective. The YOLO D instance has 19,450 parameters and an unbounded raw output, in contrast to P’s candidate convex hull. This is a comparison of readout/supervision packages, not a causal ablation of softmax alone.

L is the historical line-structured internal control. Its results are retained to show that explicit line structure has not established an all-metric benefit in this configuration. Earlier Hough, active-adaptation, and replay investigations are part of the research history, not new primary comparisons in this manuscript. No unfavorable historical outcome was deleted.

For a supported corner with target residual $\mathbf{r}_k$, P’s target mass is

$$q_{kj} = \frac{\exp[-\|d_b \mathbf{v}_j - \mathbf{r}_k\|^2 / (2\sigma^2)]}{\sum_l \exp[-\|d_b \mathbf{v}_l - \mathbf{r}_k\|^2 / (2\sigma^2)]}, \quad \sigma = 0.08d_b/17. \quad (5)$$

Cross-entropy is averaged over supported corners within each frame and then over frames with support. This smooth finite-candidate target differs from D’s direct residual L1 supervision. Its existing settings are retained, not newly optimized here.

D. Second fixed estimator: DOPE adapter

The second application freezes an existing six-stage pallet DOPE checkpoint and changes the feature adapter rather than the correction principle. The checkpoint is identified by SHA-256 in the protocol; no DOPE weight is updated by the refiner fits. We use the existing semantic-channel local-peak decoder at threshold 0.3. This implementation does not run affinity-based multi-object grouping. Its fixed box is the hull of at least three valid predicted corners, and its score is the maximum corner belief. Thus, this is an application to the declared single-pallet DOPE pipeline, not a reproduction of every original DOPE inference option or a detector-ranking experiment.

The two frozen VGG taps are post-ReLU layers 17 and 26, with 256 channels at stride 4 and 128 at stride 8. Learned projections map both to the same 16-channel evidence interface used by P and D; the shared encoded width is 24. The candidate lattice, stencil geometry, corner ordering, null candidate, losses, and displacement restoration remain as specified above. Input-channel counts change the number of learned parameters; the YOLO count is not assigned to the DOPE heads. A matched-evidence D control is trained for each DOPE seed. L and the separate PoseFix-derived comparator are historical YOLO controls and are not silently presented as new DOPE experiments.

The fixed adapter applies reflection padding of 100 pixels to native real RGB, resizes the padded canvas back to the native size, then uses the recorded shortest-side-400, multiple-of-eight raster. Synthetic source RGB already contains this reflection preparation and receives no second padding. The exact composed affine, including actual raster widths and heights, maps predictions back to the caller’s coordinates. Training retains the source preparation; DEV restores unpadded original pixels before measurement. The inference preprocessing and normalization are checked against the selected existing inference path, not claimed to be byte-identical to the original augmented DOPE training pipeline.

Predicted missing points stay missing. P and D cannot create a previously absent corner, change the fixed hull or score, or replace the center landmark. A point without the required feature support is copied rather than supplied with a target-derived value. These invariants isolate local refinement but also bound what it can repair: low coverage cannot be hidden by reporting only an accurate surviving subset. Camera intrinsics and physical dimensions remain exclusive to downstream PnP.

E. Third application: full-image ResNet18

The third estimator is a full-image pallet adaptation of Simple Baselines[8], using an ImageNet-initialized ResNet18 body, three 256-channel stride-two deconvolutions, and nine stride-four heatmaps. It uses no YOLO detection or ground-truth crop. Prepared RGB is aspect-fitted to 384×512 with centered letterboxing and ImageNet FP32 normalization. Targets are continuous Gaussians with output-grid sigma 2; the masked half-MSE retains all nine channel slots and every frame in its denominator. Decoding uses the fixed argmax/interior quarter-pixel rule and peak threshold 0.1, with explicit missingness and exact affine restoration.

The planned ResNet18 baseline is 60 epochs of 55,980 source images: 209,940 updates and 3,358,800 exposures, retaining the final 12-image batch. Adam starts at 10^{-3} with weight decay 10^{-4} and factor-ten drops at the start of epochs 40 and 50. Only calibration loss is monitored and final epoch 60 is fixed. Seed 42 and source-only photometric augmentation are prescribed. These are preparation details; actual training and outcomes have not been inserted. After the baseline is frozen, P/D read layer2/layer3 features (128/stride8 and 256/stride16) under the same fixed three-seed head budget. Each application uses independently initialized heads; no YOLO or DOPE head is a warm start.

F. Purpose-adapted prior comparator

The historical YOLO prior uses the official PoseFix ResNet152 arrangement, an aspect-corrected and 1.25-expanded predicted-box RGB crop of 384×288 , and nine Gaussian input pose maps. It adds ImageNet TF-Slim initialization, masks center loss and restores the original center. Human swap/inversion noise is replaced by actual frozen source-detector errors. Operator and training parity checks for the source-based PyTorch port, including the effective fused-BN epsilon and target-boundary handling, are detailed in the supplement.

Its declared budget is three seeds of 6,000 Adam updates at batch 16 in the exact historical P sample order, using the final checkpoint. This is a matched-exposure adaptation rather than the original 140-epoch human-pose protocol or a claim of convergence. All three matched-budget fits were completed; the learning curves are diagnostic, not a best-checkpoint selector.

TABLE I

Historical YOLO comparator contract. All added heads use zero real-pallet training examples; R0 upstream COCO-pose initialization is shared within this panel. Geometry enters PnP, not the heads. DOPE uses its own frozen checkpoint.

Method	Added inference input	Added pretraining
R0	none	none
P/D/L	R0 P3/P4, pose/box	none
Prior	RGB crop, pose/box	ImageNet classes

G. Downstream geometric measurement

The unchanged prediction-only selector uses the registered camera intrinsics and physical pallet dimensions to select its dimension hypothesis. The selected cuboid is solved using the canonical SQPnP/RefineLM path. Intrinsics and dimensions are not neural inputs to P or D. Ground truth does not choose the prediction's most favorable pose hypothesis. Current reference poses are reconstructed from manual image points and known geometry, rather than obtained from an independent physical 6D instrument. Near-square axis confusion is not newly excused by adding a symmetry minimization.

IV. Datasets and Measurement Protocol

A. Source exposure and rule selection

For historical YOLO, the source manifest contains 55,980 training records, with 55,915 usable matched training rows, and calibration/selection/heldout partitions of 1,004/1,031/1,985 records. Source images, labels, row order, and completed model choices are hash-bound. Actual RGB crops, rather than guessed images reconstructed from cached features, are used for the prior backbone. Synthetic source subsets are historically reused and are not described as universally untouched data.

The prior shares the existing synthetic-only lambda/cap grid, objective, and tie rule. Coordinate outputs do not receive a fabricated temperature sweep. Raw prior and capped-wrapper results are reported separately. Heldout results are computed after selection; DEV does not choose the checkpoint, learning rate, noise model, or wrapper. The historical model panel is frozen before any independent confirmation annotations may be opened; the DOPE extension also freezes source-selected rules before its DEV replay.

B. Fixed DOPE source budget and selection

DOPE uses the same 60,000-record synthetic RGB source manifest: 55,980 training, 1,004 calibration, 1,031 selection, and 1,985 heldout records. A usable training row needs the frozen predicted hull to meet IoU 0.5 with the source object and at least one jointly valid corner. Its usable count is estimator-specific and must not inherit the historical YOLO count. Each of three seeds trains a fresh P and a fresh D for exactly 6,000 AdamW updates of nominal batch 16. Within a seed they share the exact row sequence and frozen feature evidence, giving 96,000

nominal source exposures per head. Cross-estimator usable rows need not coincide. Frozen VGG prefixes run online; the DOPE stage heads are not rerun inside these training steps. Feature maps use the original FP16 roundtrip, with FP32 heads and no mixed-precision head training.

The learning-rate, warmup, decay, gradient clipping, and final-checkpoint rule are fixed in advance; no DEV-based early stopping or architectural sweep is allowed. A discarded two-update smoke check is a separate implementation check, not a seventh reported fit or a warm start. P temperature is selected on the calibration partition from $\{0.5, 1, 2, 4\}$ using supported-frame target cross-entropy. D has no temperature selection. Each family selects one three-seed mean rule on the selection partition from $\lambda \in \{0, 0.0625, 0.25, 1, 4\}$ and cap fractions $\{\text{none}, 0.01\}$. The fixed tie rule favors smaller λ and then the stricter cap. Missing source supervision contributes the declared failure penalty; no real image chooses a wrapper. Heldout measurements follow rule freezing and remain historically reused source data.

The baseline-training budgets are not matched: the DOPE checkpoint ending in epoch 0060 was produced by the inclusive 0–60 loop, whereas YOLO uses its selected best.pt. Shared source imagery and matched P/D head exposure do not establish equal baseline optimization, augmentation, pretraining, or compute. The supplement records this distinction alongside the source and checkpoint hashes.

C. Development population and supervision

The historical YOLO positive development panel has 319 images in 13 recorded sessions; 2689 negative images accompany the detection-preservation audit. In that historical panel there are 311 matched positive frames and 2,756 supervised points for conditional precision, versus 2,818 supervised points in the complete ground-truth denominator. Point supervision is the original label contract and is not renamed as measured physical visibility. Visible boundaries, inferred amodal corners, and unavailable landmarks require distinct interpretation in new annotation QA.

The DOPE extension reuses the identical 319 positive images and 2,818 supervised points, but reports its own matched and missing counts. Canonical conditional precision requires the fixed box match and all nine predicted points finite; a separate partial-point diagnostic cannot replace it. P and D preserve this baseline support within DOPE, without requiring equal support across estimators. There is no new DOPE negative-image/AP experiment.

The primary statistic pools supervised original-pixel Euclidean errors within each model seed, computes its median, and then averages the three seed statistics. It is neither an ensemble prediction nor the average of frame medians. Its conditional precision denominator excludes unmatched frames, which is why complete-denominator PCK is also required. For threshold t ,

$$\text{PCK}_t = \frac{\sum_{i,k} h_{ik}(t)}{N_{\text{GT}}}, \quad (6)$$

where missing predictions, unmatched top instances, and nonfinite points contribute zero hits, not zero error. Matching remains at the registered box IoU threshold. If any support differs, an intersection-only median cannot replace this failure-aware report. In the DOPE extension a conditional paired interval is not computed when its per-frame support differs; the mismatch is reported alongside full-denominator PCK and pose coverage.

Registered object sizes and the per-annotation camera intrinsics are tabulated in the supplement. They are inputs to this geometry-reference evaluation, not newly certified metrology. The canonical solver uses no distortion coefficients; independent camera calibration uncertainty is not established by these annotation records. The declared long-axis benchmark admits a half turn but not a quarter turn. In particular, the wood object’s visual symmetry under a half turn is unverified. Such conventions do not establish physical fork-entry interchangeability.

D. Uncertainty and pose metrics

We use 10,000 paired session bootstrap draws with seed 20260914, applying each draw’s frame/session multiplicities to every model seed. The P–R0 statistic compares against one baseline, not three copies treated as independent models. P–D and P–prior are secondary development contrasts. Secondary intervals are not multiplicity-adjusted and are not confirmatory superiority, equivalence, or noninferiority tests. The reused development data cannot be made independent by specifying this protocol after previous analyses.

Rotation, yaw, translation, 3D IoU, and symmetry-aware ADD AUC are evaluated under the existing geometry-reference convention. Pose coverage and 2D matching are different denominators: a solver can return a pose from an incorrectly matched top detection. Translation-error changes in centimeters can be expressed in millimeters, but a pixel reduction cannot be converted to one universal millimeter gain. No insertion trials or independently instrumented physical pose measurements are included.

E. Cost and independent confirmation

The expanded panel was measured in the same session for all comparators.

Desktop latency starts from decoded BGR input and includes preprocessing, R0, refinement, and restoration. Canonical PnP time is recorded separately and added for the whole estimation path. Loading and image decoding are excluded. The protocol uses the historical 26 images, 20 warmups per model, five balanced cyclic/reversed blocks, batch one, and four Torch intra-operation CPU threads. OpenCV/inter-operation observations were saved only for the subsequent isolated-memory phase, not the timing phase; four is not a universal process thread limit. Training and scoring do not overlap timing. Display/remote desktop activity is disclosed; no fastest repeat is selected. Prior-specific cropping and its separate backbone are included,




Fig. 3. Historical YOLO per-seed DEV median against same-session median full-path latency. Lower is preferable on both axes. Each seed is a point, not a timing-selected winner or ensemble; separate-backbone prior cost is included. Source result/runtime hashes are recorded in FIGURE_MANIFEST.json.

while P recomputes detector features instead of reading a training cache. The supplement records numerical replay checks and the recovery of memory measurement after the completed timing panel.

The DOPE timing protocol separately measures the fixed seed-1 DOPE, DOPE+D, and DOPE+P paths on the first two manifest-ordered images per session (26 images), after 20 warmups and in five balanced cycles. Its actual times, device observations, coverage, and replay checks must be reported before any cross-estimator cost claim. Historical YOLO timing is not reused as DOPE timing.

Independent confirmation is awaiting new data. The collection plan starts with 12 separate sessions and 15 prespecified frames per session, with at least 30 frames independently double annotated before blind adjudication. This is not a power guarantee. Predictions are hidden during annotation and the full panel must be frozen before evaluation. Existing FINAL-named manifests have unavailable membership and are not treated as independent observations.

V. Results

A. Completed historical YOLO comparisons

Table II summarizes the original-pixel and geometry-reference measurements. Complete per-seed tables, extra PCK thresholds, coverage, and uncertainty are supplied in the supplement. P changes the baseline pooled median from 6.616 to 5.905 pixels, a reduction of 0.711 pixels. Its paired-session reduction interval is [0.411, 1.174] pixels. PCK10 changes from 63.733% to 67.341%. These figures describe the present development panel only.

The P-D median difference is -0.525 pixels, with paired-session interval [-0.980, -0.210]. This supports a lower

TABLE II

Reused DEV: seed-mean statistics, not an ensemble. Translation is against the geometry reference. NM means not measured in this build.

Model	Median px	P90 px	PCK10 %	Trans. cm
R0	6.616	38.670	63.73	7.897
P mean	5.905	37.766	67.34	7.153
D mean	6.430	37.566	64.57	7.597
L mean	6.070	37.320	66.78	7.548
PRIOR mean	5.569	37.227	68.41	6.895




Fig. 4. P-R0 development median effects by recorded session. Negative values favor P. These are descriptive effects from the reused panel, not independent confirmation intervals. All sessions are retained.

central development residual for P than this fixed direct control. It does not prove that direct regression as a method class is inferior. D has a slightly lower observed P90, and the P-D tail interval includes zero. D also has lower measured desktop latency. There is therefore no all-metric dominance claim. The internal line control also has a lower observed P90 than P.

Figure 4 retains every recorded session. The eval_noapril session has an adverse P-R0 central effect; the other twelve observed sessions improve. Leave-one-session-out estimates remain favorable in this panel, but such deletion diagnostics are not new independent datasets. No unfavorable session is removed from the aggregate.

The seed-mean reduction in geometry-reference translation median is 7.439 mm. This is a change in error statistics against a reconstructed reference, not measured insertion improvement. It cannot establish physical metrology accuracy independently of the image-label and reference-construction uncertainty.

B. Prior-derived comparison

The prior checkpoints were fixed at the final training update; only the wrapper rule was selected on synthetic data. The selected wrapper has lambda 1.000 and cap fraction 0.01. Raw and wrapped outputs are separate supplementary rows. All wrapped prior seeds retain the baseline matched supervision support, with no nonfinite supervised coordinates. The P-prior paired median difference is 0.336 px, with session interval [0.063, 0.585] px.

TABLE III

Desktop median latency in ms, representative seed1.
Training/scoring excluded from timing; all repeats retained.

Model	BGR-to-2D	PnP	Total
R0	9.733	1.422	11.217
P1	13.661	1.382	15.352
D1	13.231	1.416	14.800
L1	18.226	1.399	19.736
PRIOR1	27.537	1.470	29.020

This interval favors the prior, so P is not the most accurate comparator under this central-error criterion. A negative contrast favors P; a positive contrast favors the prior. The complete-denominator PCK, tail and pose rows must be considered alongside the median. Six thousand updates per seed do not certify saturation or reproduce the original 140-epoch protocol.

C. Second-estimator development evidence

DOPE source training/evaluation results have not been inserted. No DOPE accuracy, coverage, pose benefit, parameter count or measured latency is claimed in this draft. DOPE measurement status: pending. The reserved result table will contain baseline DOPE and three-seed D/P statistics, with matched counts, all-GT PCK and pose coverage. No zero or historical YOLO value represents an unmeasured DOPE result. The supplement reports the full per-seed table, paired intervals, source-heldout statistics and selected wrappers. Within DOPE, central residuals, complete-denominator accuracy, tails, and pose coverage are separate outcomes. A lower conditional median alone cannot establish a net sensing improvement when missing predictions remain. Across estimators, unequal baseline errors, boxes, features, and upstream training prevent a causal attribution to backbone identity. Any favorable DOPE effect supports this fixed application only; a null or adverse result remains part of the same experiment.

D. Third-estimator development evidence

ResNet18 baseline completion and before/after measurements remain pending. ResNet18 measurement status: pending. The reserved result table will contain baseline ResNet18 and three-seed D/P statistics, with matched counts, all-GT PCK and pose coverage. No zero or historical YOLO value represents an unmeasured ResNet18 result. The supplement retains all seven arms, paired intervals, source evidence, the complete baseline learning curve and actual timing. Cross-estimator comparisons remain descriptive because their baseline training and matched supports differ.

E. Efficiency and accuracy–cost tradeoffs

The expanded panel was measured in the same session for all comparators.

P’s selected seed1 is the prespecified deployment representative; it must not be confused with the seed mean. Its

desktop BGR-to-2D and PnP-inclusive median times are 13.661 and 15.352 ms, compared with 9.733 and 11.217 ms for R0 in the cited timing session. Allocator memory is reported separately from whole-device memory, which also includes the display and other software. The tables do not extrapolate to Jetson, report unmeasured energy savings, or infer operational throughput from an isolated cached tensor.

The separate DOPE and ResNet18 seed-1 timing panels, including all pose failures, full-path timings and coordinate replay checks, are reported in the supplement. Their baselines and feature extraction paths differ; historical YOLO timing is not a substitute for either measurement.

F. Independent confirmation

Not yet measured: no independently reviewed capture/annotation panel is available. The proposed collection count is a plan, not a completed experiment. Development observations are not relabeled as confirmation results.

VI. Limitations and Interpretation

First, the primary development improvement is conditional on matched supervision and on a historically selected model family. Complete-denominator PCK limits one failure mode of optimistic reporting, but does not turn reused DEV into a confirmation set. The session intervals condition on only thirteen observed sessions and three trained seeds; they do not quantify every deployment shift.

Second, the comparison contracts differ in useful ways that must remain visible. P reads internal R0 features; the PoseFix-derived prior computes a separate RGB representation and adds ImageNet pretraining. D has a different output support and loss. Equal source exposure is not equal compute or equal optimization difficulty. None of these comparisons identifies a universal architectural optimum. Bounded prior learning curves diagnose progress but cannot certify full convergence.

Third, current pose errors are measured against a geometry-reconstructed reference. Annotation ambiguity, calibration, object dimensions, and reference construction can affect interpretation. Independently blinded duplicate labels and reviewed capture lineage are required for new confirmation; physical 6D accuracy would require additional metrology. Neither preserved detections nor improved median error guarantees harmless tails, correct pallet axes, or safe handling.

Finally, each application is tied to a declared frozen feature interface and a desktop environment. The DOPE application requires its own channel/stride adapter, missing-point policy, new head fits, and decoder audit; it is not plug-and-play evidence for an arbitrary estimator. Reusing DEV for the second application does not create independent confirmation. Embedded export parity and Jetson latency are not measured. The work provides an auditable sensing experiment, not a demonstrated control

system. The author team must assess whether the remaining independent evidence and scientific contribution are sufficient for submission; compilation and formatting checks are not an acceptance decision.

VII. Conclusion

The completed YOLO application of a fixed feature-reusing point refiner improves the central landmark residual on the reused pallet development panel with a measured desktop cost. The preserved detection interface, failure-aware denominators, internal controls, and source-derived prior comparison make the scope of this claim explicit. The source-derived prior comparison quantifies a fixed-budget tradeoff, not converged-method superiority. Independent real-session confirmation remains required. The DOPE application remains unresolved; the historical YOLO result is not evidence of its outcome. The required ResNet18 application remains unresolved; neither a completed interface nor results from another estimator establish its outcome. Further claims about independent generalization, physical sensing accuracy, or operational success require the corresponding evidence rather than a more favorable presentation of development numbers.

Data, Code, and Author Declarations

Scripts, small numerical artifacts, coordinate checks, source/weight hashes, and reproduction commands are assembled in the companion repository paths identified in the supplement. Raw images, videos, and large external/model weights are not publicly redistributed without permission. The source ImageNet checkpoint retrieval procedure is documented separately. Author names, affiliations, ORCID identifiers, funding, conflicts, related prior publications, AI-use disclosure, and submission consent remain for the actual author team to confirm. This internal draft has not been submitted.

References

- [1] M. Knitt, J. Schyga, A. Adamanov, J. Hinckeldeyn, and J. Kreutzfeldt, “Estimating the pose of a euro pallet with an RGB camera based on synthetic training data,” *Logistics Journal: Proceedings*, no. 18, 2022.
- [2] J. Tremblay, T. To, B. Sundaralingam, Y. Xiang, D. Fox, and S. Birchfield, “Deep object pose estimation for semantic robotic grasping of household objects,” in *Conference on Robot Learning*, 2018. [Online]. Available: <https://arxiv.org/abs/1809.10790>
- [3] A. Miura, H. Uchiyama, M. Yamaguchi, N. Kai, T. Shirosima, and H. Saito, “Real-time 6DoF pallet pose estimation with monocular metric depth,” in *International Conference on Computer Vision Theory and Applications*, vol. 2, 2026, pp. 149–161. [Online]. Available: <https://www.scitepress.org/PublishedPapers/2026/146268/>
- [4] G. Moon, J. Y. Chang, and K. M. Lee, “PoseFix: Model-agnostic general human pose refinement network,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019, pp. 7773–7781.
- [5] X. Sun, B. Xiao, F. Wei, S. Liang, and Y. Wei, “Integral human pose regression,” in *European Conference on Computer Vision*, 2018, pp. 536–553.
- [6] P. Castro and T.-K. Kim, “CRT-6D: Fast 6D object pose estimation with cascaded refinement transformers,” in *IEEE/CVF Winter Conference on Applications of Computer Vision*, 2023, pp. 5735–5744. [Online]. Available: <https://arxiv.org/abs/2210.11718>

- [7] J. O. Zhang, A. Sax, A. Zamir, L. Guibas, and J. Malik, “Side-tuning: A baseline for network adaptation via additive side networks,” in *European Conference on Computer Vision*, 2020, pp. 698–714. [Online]. Available: https://www.ecva.net/papers/eccv_2020/papers_ECCV/html/1104_ECCV_2020_paper.php
- [8] B. Xiao, H. Wu, and Y. Wei, “Simple baselines for human pose estimation and tracking,” in *European Conference on Computer Vision*, 2018, pp. 466–481. [Online]. Available: https://www.ecva.net/papers/eccv_2018/papers_ECCV/html/Bin_Xiao_Simple_Baselines_for_ECCV_2018_paper.php